

```
import numpy as np

import pandas as pd

df = pd.read_csv(r"C:\Users\310623205040\Downloads\diabetes_dataset.csv")

df.head()

df.shape

df.dtypes

df.info()

df.describe()

sr = pd.Series(df['class'])

pos = sr.value_counts()

print(pos)

df['Age'].mean()

df['Age'].median()

df['Age'].mode()

df['Age'].std()

df['Age'].var()

df['Age'].skew()

df['Age'].kurt()

import seaborn as sns

import matplotlib.pyplot as plt

sns.histplot(df['Age'], kde=True)

plt.show()

import matplotlib.pyplot as plt
```

```

import pandas as pd

import numpy

from sklearn import datasets, linear model

from sklearn.metrics import mean_squared_error, r2_score

from sklearn, linear model import LogisticRegression

#To calculate accuracy measures and confusion matrix from sklearn import metrics

from sklearn.model_selection import train_test_split

diabetes_X, diabetes_y = datasets.load_diabetes(return_X_y=True)

diabetes_X_train, diabetes_X_test, diabetes_y_train, diabetes_y_test = train_test_split(
    diabetes_X, diabetes_y, test_size=0.2, random_state=0)

#Split the targets into training/testing sets

diabetes_X_train, diabetes_X_test, diabetes_y_train, diabetes_y_test = train_test_split(
    diabetes_X, diabetes_y, test_size=0.2, random_state=0)

#Create Linear regression object

regr = linear_model.LinearRegression()

#Train the model using the training sets
regr.fit(diabetes_X_train, diabetes_y_train)

#Make predictions using the testing set

diabetes_X_test, diabetes_y_test, diabetes_pred = train_test_split(
    diabetes_X, diabetes_y, test_size=0.2, random_state=0)

#Create Logistic regression object
Logistic model = LogisticRegression()

Logistic model.fit(diabetes_X_train, diabetes_y_train)

#The coefficients

print('Coefficients: \n', regr.coef_)

#The mean squared error

```

```
print('Mean squared error: %.2f mean squared error(diabetes y test, diabetes_y_pred))
```

The coefficient of determination 1 is perfect prediction

```
print('Coefficient of determination: 3.21 r2 score(diabetes_y_test, diabetes_y_pred))
```

```
y_predict Logistic_model.predict(diabetes_X_train)
```

```
#print("y predict/hat", y_predict)
```

```
y_predict
```

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
diabetes_df = pd.read_csv('https://raw.githubusercontent.com/umishrabi/MachineLearning/master/Datasets/diabetes.csv')
diabetes_df.head()
```

```
X = diabetes_df.drop('Outcome', axis=1)
```

```
Y = diabetes_df['Outcome']
```

```
from sklearn.preprocessing import MinMaxScaler
```

```
scaler = MinMaxScaler(feature_range=(0,1))
```

```
scaled_data = scaler.fit_transform(X)
```

```
scaled_data
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, Y_train, Y_test = train_test_split(scaled_data, Y, test_size=0.2, random_state=0)
```

```
#Linear Regression
```

```
from sklearn.linear_model import LinearRegression
```

```
lin_reg = LinearRegression()
```

```
Initialize Linear Regression model (REMOVE normalize)
```

```
lin_reg = LinearRegression(copy_X=True, fit_intercept=True, n_jobs=-1)
```

```
Fit the model
```

```
lin_reg.fit(X_train, Y_train)
```

```
#Get score
```

```
r2lin_reg.score(X_test, Y_test)

print (f"Linear Regression Score: (r2:.4f)")

from sklearn.metrics import mean_squared_error, r2_score

predictions = lin_reg.predict(X_test)

predictions

mean_squared_error(Y_test, predictions)

r2_score (Y_test, predictions)
```

```

import pandas as pd

import numpy as np

#Read CSV files

data_1= pd.read_csv(r"C:\Users\310623205040\Downloads\car1.csv")

data_2= pd.read_csv(r"C:\Users\310623205040\Downloads\car2.csv")

Convert to Datoframe

df1= pd.DataFrame(data_1)

df2= pd.DataFrame(data_2)

Ensure both datoframes have the amount and amount columns

if "Amount" in df1.columns and "Amount1" in df2.columns:

Assign amount1 column from df2 to df 1

df1["Amount1"]=df2["Amount"]

Check if prices match

df1["prices_match"]= np.where(df1["Amount"]==df1["Amount1"], True, False)

Compute price difference

df1["price_diff"]= np.where(df1["Amount"]!=df1["Amount1"], df1["Amount"]-df1["Amount1"], 0)

Print Dataframe

print(df1)

else:

print("Error: Required columns ('Amount' in df1 and 'Amount 1' in df2) are missing.")

```

```

import numpy as np

import matplotlib.pyplot as plt

5/11

from scipy.stats import norm

%matplotlib inline.

#define constants

in 998.8

sigma 73.10

x1 = 900

x2= 1100

#calculate the z-transform

z1(x1mu) / sigma

z2 = (x2mu) / sigma

x= np.arange(z1, z2, 0.001) range of x in spec

x_all= np.arange(-10, 10, 8.001) # entire range of x, both in and out of spec

#mean 8, stddev 1, since Z-transform was calculated

y norm.pdf(x,0,1)

y2 norm.pdf(x_all,0,1)

#build the plot

fig, ax plt.subplots(figsize=(9,6))

plt.style.use('fivethirtyeight')

ax.plot(x_all,y2)

ax.fill_between (x,y,, alpha=0.3, color='b')

ax.fill_between(x_all,y2,0, alpha=0.1)

```

```
ax.set_xlim([-4,4])  
  
ax.set_xlabel('# of Standard Deviations Outside the Mean')  
  
ax.set_yticklabels([])  
  
ax.set_title('Normal Gaussian Curve')  
  
plt.savefig('normal_curve.png', dpi=72, bbox_inches='tight')  
  
plt.show()
```



```

import pandas as pd

iris = pd.read_csv(r"C:\Users\310623205040\Downloads\iris_data.csv")

print(iris.head())

import pandas as pd

import matplotlib.pyplot as plt

from seaborn import load_dataset

iris = load_dataset("iris")

iris.plot(kind='density', subplots=True, layout=(4,4), sharex=False, figsize=(10, 8))

plt.show()

import matplotlib

import matplotlib.pyplot as plt

import seaborn as sns

data = sns.load_dataset("iris")

data.head()

silky_data = data[data.species == 'setosa']

virginica_data = data[data.species == 'virginica']

plt.title("Flowers (Setosa & Virginica)")

sns.kdeplot(x=silky_data.sepal_length, y=silky_data.sepal_width, shade=True, cmap='inferno', shade_lowest=False)
sns.kdeplot(x=virginica_data.sepal_length, y=virginica_data.sepal_width, shade=True, cmap='inferno', shade_lowest=False);

```

```

import pandas as pd

#Load the dataset

iris = pd.read_csv(r"C:\Users\310623205040\Downloads\iris_data.csv")

#Exclude non-numeric columns

iris_numeric = iris.select_dtypes(include=['number'])

#Compute correlation

correlation_matrix = iris_numeric.corr()

print(correlation_matrix)

import pandas as pd

import seaborn as sb

import matplotlib.pyplot as plt

#Load dataset

iris = pd.read_csv(r"C:\Users\310623205040\Downloads\iris_data.csv")

#Drop the categorical column 'species'

iris_numeric = iris.select_dtypes(include=['number'])

#Create the heatmap

fig, ax = plt.subplots(figsize=(8,6)) # Set figure size

plt.title("Iris Correlation Plot")

sb.heatmap(iris_numeric.corr(), annot=True, cmap='coolwarm', ax=ax)

plt.show()

import pandas as pd

iris =

pd.read_csv(r"C:\Users\310623205040\Downloads\iris_data.c
print(iris.head())
sv")

```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

Load dataset

```
iris = pd.read_csv("C:\Users\310623205040\Downloads\iris_data.csv")
```

Print actual column names to verify

```
print(iris.columns)
```

```
#Standardize column names
```

```
iris.columns = iris.columns.str.strip().str.lower()
```

Define color mapping for species

```
colors = {'setosa': 'red', 'versicolor': 'blue', 'virginica': 'green'}
```

Create scatter plot

```
fig, ax = plt.subplots()
```

```
for i in range(len(iris)):
```

```
    ax.scatter(iris["sepal length"][i], iris["sepal width"][i], color=colors[iris["species"][i]])
```

```
plt.xlabel("Sepal Length")
```

```
plt.ylabel("Sepal Width")
```

```
plt.title("Iris Dataset Scatter Plot")
```

```
plt.show()
```

```
import pandas as pd

iris =

pd.read_csv(r"C:\Users\310623205040\Downloads\iris_data.csv")
print(iris.head())

fig=plt.figure(figsize = (15,20))

ax=fig.gca()

iris.hist(ax = ax)

plt.show()
```

```
import pandas as pd

import numpy as np

import matplotlib.pyplot as plt # To visualize

from mpl_toolkits.mplot3d import Axes3D

data = pd.read_csv(r"C:\Users\310623205040\Downloads\pima_diabetes.csv")

data.head()

fig = plt.figure(figsize=(4,4))

ax = fig.add_subplot(111, projection='3')

fig = plt.figure()

ax = fig.add_subplot(111, projection='3d')

x = data['Age'].values

y = data['Glucose'].values

z = data['Outcome'].values

ax.set_xlabel("Age (Year)")

ax.set_ylabel("Glucose (Reading)")

ax.set_zlabel("Outcome (8 or 1)")

ax.scatter(x, y, z, c='r', marker='o')

plt.show()
```

```
from mpl_toolkits.basemap import Basemap

import matplotlib.pyplot as plt

fig plt.figure(figsize (12,12))

m = Basemap()

m.drawcoastlines (linewidth=1.0, linestyle='dashed', color='red')

plt.title("Coastlines", fontsize=20)

plt.show()

fig plt.figure(figsize (12,12))

m = Basemap()

m.drawcoastlines (linewidth=1.0, linestyle='solid', color='black')

m.drawcountries (linewidth=1.0, linestyle='solid', color='k')

plt.title("Country boundaries", fontsize=20)

plt.show()

fig plt.figure(figsize (12,12))

m = Basemap()

m.drawcoastlines (linewidth=1.0, linestyle' solid', color='black')

m.drawcountries (linewidth=1.0, linestyle'solid', color='k')

m.drawrivers (linewidth=0.5, linestyle'solid', color='#0000ff")

plt.title("Major rivers", fontsize=20)

plt.show()

fig plt.figure(figsize = (12,12))

m = Basemap()
```

```
m.drawcoastlines (linewidth=1.0, linestyle='solid', color='black')
```

```
m.drawcountries (linewidth=1.0, linestyle='solid', color='k')
```

```
m.fillcontinents (color='coral', lake_color='aqua', alpha=0.9)
```

```
plt.title("Color filled continents", fontsize=20)
```

```
plt.show()
```